Volume 03 - Book 03.03 AI Society

Version v5 draft

README.md

Book 03.03 - AI Society

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-README

Purpose

Define the AI Society blueprint package and its subsystem decomposition as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for the AI Society blueprint package and its subsystem decomposition. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for the AI Society blueprint package and its subsystem decomposition.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for the AI Society blueprint package and its subsystem decomposition.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.

- Human review and governance checkpoint interfaces.

Events

- AiSocietyDefined
- AiSocietyAssigned
- AiSocietyStateChanged
- AiSocietyReviewRequested
- AiSocietyGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- AI Society has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006
- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03 Index

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-INDEX

Purpose

Define the canonical reading order and document IDs for AI Society as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for the canonical reading order and document IDs for AI Society. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for the canonical reading order and document IDs for AI Society.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for the canonical reading order and document IDs for AI Society.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- Book0303IndexDefined

- Book0303IndexAssigned
- Book0303IndexStateChanged
- Book0303IndexReviewRequested
- Book0303IndexGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Book 03.03 Index has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006
- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005

- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.01 - AI Society Overview

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-001

Purpose

Define the governed society of AI chiefs, managers, swarms, agents, skills, capabilities, memory, communication, planning, execution, review, validation, optimization, DNA, and governance as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for the governed society of AI chiefs, managers, swarms, agents, skills, capabilities, memory, communication, planning, execution, review, validation, optimization, DNA, and governance. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for the governed society of AI chiefs, managers, swarms, agents, skills, capabilities, memory, communication, planning, execution, review, validation, optimization, DNA, and governance.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for the governed society of AI chiefs, managers, swarms, agents, skills, capabilities, memory, communication, planning, execution, review, validation, optimization, DNA, and governance.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.

- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AiSocietyOverviewDefined
- AiSocietyOverviewAssigned
- AiSocietyOverviewStateChanged
- AiSocietyOverviewReviewRequested
- AiSocietyOverviewGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- AI Society Overview has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006
- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.02 - Chief

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-002

Purpose

Define the AI Chief role responsible for strategic orchestration, mission acceptance, authority boundaries, escalation, and executive-grade oversight as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for the AI Chief role responsible for strategic orchestration, mission acceptance, authority boundaries, escalation, and executive-grade oversight. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for the AI Chief role responsible for strategic orchestration, mission acceptance, authority boundaries, escalation, and executive-grade oversight.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for the AI Chief role responsible for strategic orchestration, mission acceptance, authority boundaries, escalation, and executive-grade oversight.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- ChiefDefined
- ChiefAssigned
- ChiefStateChanged
- ChiefReviewRequested
- ChiefGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Chief has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.03 - Mission Manager

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-003

Purpose

Define mission decomposition, assignment, state tracking, dependency management, and completion evidence for AI-led work as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for mission decomposition, assignment, state tracking, dependency management, and completion evidence for AI-led work. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for mission decomposition, assignment, state tracking, dependency management, and completion evidence for AI-led work.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for mission decomposition, assignment, state tracking, dependency management, and completion evidence for AI-led work.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- MissionManagerDefined
- MissionManagerAssigned
- MissionManagerStateChanged
- MissionManagerReviewRequested
- MissionManagerGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Mission Manager has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.04 - Department Managers

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-004

Purpose

Define AI department managers that coordinate specialized agent groups against enterprise departments and missions as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for AI department managers that coordinate specialized agent groups against enterprise departments and missions. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for AI department managers that coordinate specialized agent groups against enterprise departments and missions.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for AI department managers that coordinate specialized agent groups against enterprise departments and missions.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- DepartmentManagersDefined
- DepartmentManagersAssigned
- DepartmentManagersStateChanged
- DepartmentManagersReviewRequested
- DepartmentManagersGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Department Managers has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.05 - Team Leaders

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-005

Purpose

Define team leader agents that supervise task groups, local planning, review routing, and team-level accountability as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for team leader agents that supervise task groups, local planning, review routing, and team-level accountability. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for team leader agents that supervise task groups, local planning, review routing, and team-level accountability.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for team leader agents that supervise task groups, local planning, review routing, and team-level accountability.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- TeamLeadersDefined
- TeamLeadersAssigned
- TeamLeadersStateChanged
- TeamLeadersReviewRequested
- TeamLeadersGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Team Leaders has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.06 - Swarms

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-006

Purpose

Define dynamic swarms for parallel reasoning, critique, consensus, recovery, and coordinated execution as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for dynamic swarms for parallel reasoning, critique, consensus, recovery, and coordinated execution. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for dynamic swarms for parallel reasoning, critique, consensus, recovery, and coordinated execution.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for dynamic swarms for parallel reasoning, critique, consensus, recovery, and coordinated execution.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- SwarmsDefined
- SwarmsAssigned
- SwarmsStateChanged
- SwarmsReviewRequested
- SwarmsGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Swarms has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.07 - Agents

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-007

Purpose

Define individual agent identity, role contracts, lifecycle, permissions, tools, outputs, and failure behavior as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for individual agent identity, role contracts, lifecycle, permissions, tools, outputs, and failure behavior. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for individual agent identity, role contracts, lifecycle, permissions, tools, outputs, and failure behavior.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for individual agent identity, role contracts, lifecycle, permissions, tools, outputs, and failure behavior.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentsDefined
- AgentsAssigned
- AgentsStateChanged
- AgentsReviewRequested
- AgentsGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agents has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.08 - Skills

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-008

Purpose

Define agent skills as bounded reusable knowledge and action patterns with inputs, constraints, evaluation, and versioning as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for agent skills as bounded reusable knowledge and action patterns with inputs, constraints, evaluation, and versioning. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for agent skills as bounded reusable knowledge and action patterns with inputs, constraints, evaluation, and versioning.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for agent skills as bounded reusable knowledge and action patterns with inputs, constraints, evaluation, and versioning.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- SkillsDefined
- SkillsAssigned
- SkillsStateChanged
- SkillsReviewRequested
- SkillsGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Skills has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.09 - Capabilities

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-009

Purpose

Define AI Society capabilities that package repeatable outcomes across agents, skills, tools, models, and governance gates as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for AI Society capabilities that package repeatable outcomes across agents, skills, tools, models, and governance gates. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for AI Society capabilities that package repeatable outcomes across agents, skills, tools, models, and governance gates.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for AI Society capabilities that package repeatable outcomes across agents, skills, tools, models, and governance gates.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- CapabilitiesDefined
- CapabilitiesAssigned
- CapabilitiesStateChanged
- CapabilitiesReviewRequested
- CapabilitiesGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Capabilities has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.10 - Agent Memory

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-010

Purpose

Define agent memory scopes, provenance, retention, correction, recall, and separation between temporary context and governed knowledge as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for agent memory scopes, provenance, retention, correction, recall, and separation between temporary context and governed knowledge. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for agent memory scopes, provenance, retention, correction, recall, and separation between temporary context and governed knowledge.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for agent memory scopes, provenance, retention, correction, recall, and separation between temporary context and governed knowledge.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentMemoryDefined
- AgentMemoryAssigned
- AgentMemoryStateChanged
- AgentMemoryReviewRequested
- AgentMemoryGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent Memory has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.11 - Agent Communication

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-011

Purpose

Define message contracts, handoffs, broadcasts, critiques, escalation messages, and communication auditability as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for message contracts, handoffs, broadcasts, critiques, escalation messages, and communication auditability. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for message contracts, handoffs, broadcasts, critiques, escalation messages, and communication auditability.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for message contracts, handoffs, broadcasts, critiques, escalation messages, and communication auditability.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentCommunicationDefined
- AgentCommunicationAssigned
- AgentCommunicationStateChanged
- AgentCommunicationReviewRequested
- AgentCommunicationGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent Communication has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

book-03-03-12-agent-planning.md

Book 03.03.12 - Agent Planning

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-012

Purpose

Define planning primitives for goals, tasks, constraints, dependencies, risks, options, and plan approval as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for planning primitives for goals, tasks, constraints, dependencies, risks, options, and plan approval. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for planning primitives for goals, tasks, constraints, dependencies, risks, options, and plan approval.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for planning primitives for goals, tasks, constraints, dependencies, risks, options, and plan approval.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentPlanningDefined
- AgentPlanningAssigned
- AgentPlanningStateChanged
- AgentPlanningReviewRequested
- AgentPlanningGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent Planning has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.13 - Agent Execution

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-013

Purpose

Define execution control for agent actions, tool use, runtime state, retries, rollback, and human-in-the-loop gates as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for execution control for agent actions, tool use, runtime state, retries, rollback, and human-in-the-loop gates. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for execution control for agent actions, tool use, runtime state, retries, rollback, and human-in-the-loop gates.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for execution control for agent actions, tool use, runtime state, retries, rollback, and human-in-the-loop gates.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentExecutionDefined
- AgentExecutionAssigned
- AgentExecutionStateChanged
- AgentExecutionReviewRequested
- AgentExecutionGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent Execution has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.14 - Agent Review

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-014

Purpose

Define review loops for agent outputs, peer critique, human approval, revision, and decision evidence as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for review loops for agent outputs, peer critique, human approval, revision, and decision evidence. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for review loops for agent outputs, peer critique, human approval, revision, and decision evidence.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for review loops for agent outputs, peer critique, human approval, revision, and decision evidence.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentReviewDefined
- AgentReviewAssigned
- AgentReviewStateChanged
- AgentReviewReviewRequested
- AgentReviewGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent Review has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.15 - Agent Validation

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-015

Purpose

Define validation of agent claims, outputs, constraints, tool results, safety conditions, and acceptance evidence as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for validation of agent claims, outputs, constraints, tool results, safety conditions, and acceptance evidence. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for validation of agent claims, outputs, constraints, tool results, safety conditions, and acceptance evidence.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for validation of agent claims, outputs, constraints, tool results, safety conditions, and acceptance evidence.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentValidationDefined
- AgentValidationAssigned
- AgentValidationStateChanged
- AgentValidationReviewRequested
- AgentValidationGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent Validation has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.16 - Agent Optimization

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-016

Purpose

Define measured improvement of prompts, skills, plans, routing, model selection, and swarm patterns without uncontrolled self-modification as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for measured improvement of prompts, skills, plans, routing, model selection, and swarm patterns without uncontrolled self-modification. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for measured improvement of prompts, skills, plans, routing, model selection, and swarm patterns without uncontrolled self-modification.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for measured improvement of prompts, skills, plans, routing, model selection, and swarm patterns without uncontrolled self-modification.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentOptimizationDefined
- AgentOptimizationAssigned
- AgentOptimizationStateChanged
- AgentOptimizationReviewRequested
- AgentOptimizationGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent Optimization has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.17 - Agent DNA

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-017

Purpose

Define agent DNA packages that encode role identity, behavioral constraints, preferences, operating patterns, and evolution metadata as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for agent DNA packages that encode role identity, behavioral constraints, preferences, operating patterns, and evolution metadata. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for agent DNA packages that encode role identity, behavioral constraints, preferences, operating patterns, and evolution metadata.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for agent DNA packages that encode role identity, behavioral constraints, preferences, operating patterns, and evolution metadata.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentDnaDefined
- AgentDnaAssigned
- AgentDnaStateChanged
- AgentDnaReviewRequested
- AgentDnaGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent DNA has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

book-03-03-18-agent-governance.md

Book 03.03.18 - Agent Governance

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-018

Purpose

Define governance controls for AI authority, policies, protected actions, exceptions, approvals, and audit trails as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for governance controls for AI authority, policies, protected actions, exceptions, approvals, and audit trails. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for governance controls for AI authority, policies, protected actions, exceptions, approvals, and audit trails.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for governance controls for AI authority, policies, protected actions, exceptions, approvals, and audit trails.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AgentGovernanceDefined
- AgentGovernanceAssigned
- AgentGovernanceStateChanged
- AgentGovernanceReviewRequested
- AgentGovernanceGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- Agent Governance has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

Book 03.03.19 - AI Society Testing

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-019

Purpose

Define testing strategy for AI Society behavior, coordination, safety, memory, execution, validation, and governance as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for testing strategy for AI Society behavior, coordination, safety, memory, execution, validation, and governance. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for testing strategy for AI Society behavior, coordination, safety, memory, execution, validation, and governance.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for testing strategy for AI Society behavior, coordination, safety, memory, execution, validation, and governance.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AiSocietyTestingDefined
- AiSocietyTestingAssigned
- AiSocietyTestingStateChanged
- AiSocietyTestingReviewRequested
- AiSocietyTestingGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- AI Society Testing has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.

book-03-03-20-ai-society-roadmap.md

Book 03.03.20 - AI Society Roadmap

Version: v5 draft

Status: Draft

Document ID: MAL-V03-B03-03-020

Purpose

Define the implementation-facing roadmap for moving AI Society from blueprint to specifications and governed runtime behavior as part of the Mariam AI Society blueprint.

Scope

This document covers architecture-level responsibilities, contracts, data, events, controls, metrics, tests, and acceptance rules for the implementation-facing roadmap for moving AI Society from blueprint to specifications and governed runtime behavior. It does not implement system code or grant autonomous authority beyond documented governance.

Responsibilities

- Establish the canonical boundary for the implementation-facing roadmap for moving AI Society from blueprint to specifications and governed runtime behavior.
- Coordinate with Enterprise Core identity, event, storage, security, and observability services.
- Preserve human review, traceability, and rollback paths for high-impact AI outcomes.
- Produce implementation-ready constraints for future specifications.

Inputs

- Enterprise Constitution AI, swarm, security, governance, quality, and engineering principles.
- Master Roadmap Phase 03 AI Society requirements.
- Enterprise Core runtime, identity access, event bus, storage, and security contracts.
- Enterprise Organization missions, departments, escalation paths, and approval lines.

Outputs

- Blueprint contract for the implementation-facing roadmap for moving AI Society from blueprint to specifications and governed runtime behavior.
- Governed AI Society records, events, and acceptance expectations.
- Traceable inputs for Volume 06 specifications and Volume 11 AI Society expansion.
- Reviewable evidence for future implementation gates.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, service container, storage, and observability interfaces.
- Enterprise Organization decision flow, escalation model, approval lines, and department mission interfaces.
- Future capability system, knowledge system, model ecosystem, and workflow runtime interfaces.
- Human review and governance checkpoint interfaces.

Events

- AiSocietyRoadmapDefined
- AiSocietyRoadmapAssigned
- AiSocietyRoadmapStateChanged
- AiSocietyRoadmapReviewRequested
- AiSocietyRoadmapGovernanceBlocked

Storage

- AI Society registry records for the relevant chief, manager, swarm, agent, skill, capability, memory, or governance object.
- Mission, plan, execution, review, validation, and optimization journals.
- Audit metadata linking agent behavior to user intent, source context, model choice, and approval status.
- Retention metadata separating temporary reasoning context from approved knowledge or DNA.

Security

- Apply least privilege to agents, tools, models, runtimes, memories, and protected actions.
- Deny autonomous mutation of authoritative enterprise records unless an approved workflow grants authority.
- Classify prompts, memory, context, outputs, and event payloads before storage or broadcast.
- Audit privileged agent decisions, denied actions, escalations, model changes, and governance exceptions.

Metrics

- Task completion quality and acceptance rate.
- Review rejection rate and revision count.
- Escalation rate, blocked-action count, and governance exception count.
- Latency, cost, model usage, memory recall precision, and validation pass rate.

Tests

- Contract tests for AI Society interfaces and event envelopes.
- Safety tests for protected actions, role boundaries, and tenant isolation.
- Evaluation tests for output quality, validation, memory recall, planning, and execution.
- Failure-mode tests for escalation, rollback, retry limits, and human review routing.

Acceptance Criteria

- AI Society Roadmap has explicit authority boundaries and review paths.
- Events, storage records, metrics, and tests are defined well enough for downstream specifications.
- Security and governance controls prevent hidden autonomous execution.
- Traceability links this blueprint to Volume 00, Volume 02, Book 03.01, and Book 03.02.

Traceability

- MAL-V00-B006

- MAL-V00-B008
- MAL-V00-B011
- MAL-V02-B005
- MAL-V03-B03-01-004
- MAL-V03-B03-02-006

Version History

Version	Date	Status	Notes
v5 draft	2026-06-29	Draft	Initial AI Society blueprint content for Mariam Architecture Library v5.