

Volume 03 - Book 03.20 Autonomous Enterprise

Version v22 draft

README.md

Book 03.20 - Autonomous Enterprise

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-README

Purpose

Define the Autonomous Enterprise blueprint package and its subsystem decomposition as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for the Autonomous Enterprise blueprint package and its subsystem decomposition. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for the Autonomous Enterprise blueprint package and its subsystem decomposition.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for the Autonomous Enterprise blueprint package and its subsystem decomposition.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.

- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomousEnterpriseDefined
- AutonomousEnterpriseRegistered
- AutonomousEnterpriseStateChanged
- AutonomousEnterpriseReviewRequested
- AutonomousEnterpriseGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.

- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.
- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomous Enterprise has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20 - Autonomous Enterprise

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-INDEX

Purpose

Define the Autonomous Enterprise blueprint package and its subsystem decomposition as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for the Autonomous Enterprise blueprint package and its subsystem decomposition. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for the Autonomous Enterprise blueprint package and its subsystem decomposition.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for the Autonomous Enterprise blueprint package and its subsystem decomposition.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomousEnterpriseDefined
- AutonomousEnterpriseRegistered
- AutonomousEnterpriseStateChanged
- AutonomousEnterpriseReviewRequested
- AutonomousEnterpriseGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomous Enterprise has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.01 - Autonomous Enterprise Overview

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-001

Purpose

Define Autonomous Enterprise Overview within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomous Enterprise Overview within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation- first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomous Enterprise Overview within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomous Enterprise Overview within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomousEnterpriseOverviewDefined
- AutonomousEnterpriseOverviewRegistered
- AutonomousEnterpriseOverviewStateChanged
- AutonomousEnterpriseOverviewReviewRequested
- AutonomousEnterpriseOverviewGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomous Enterprise Overview has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.02 - Autonomy Levels

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-002

Purpose

Define Autonomy Levels within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomy Levels within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomy Levels within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomy Levels within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.

- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomyLevelsDefined
- AutonomyLevelsRegistered
- AutonomyLevelsStateChanged
- AutonomyLevelsReviewRequested
- AutonomyLevelsGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.
- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.

- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomy Levels has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.03 - Autonomous Planning

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-003

Purpose

Define Autonomous Planning within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomous Planning within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomous Planning within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomous Planning within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomousPlanningDefined
- AutonomousPlanningRegistered
- AutonomousPlanningStateChanged
- AutonomousPlanningReviewRequested
- AutonomousPlanningGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomous Planning has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.04 - Supervised Execution

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-004

Purpose

Define Supervised Execution within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Supervised Execution within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Supervised Execution within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Supervised Execution within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- SupervisedExecutionDefined
- SupervisedExecutionRegistered
- SupervisedExecutionStateChanged
- SupervisedExecutionReviewRequested
- SupervisedExecutionGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Supervised Execution has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.05 - Enterprise Simulation

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-005

Purpose

Define Enterprise Simulation within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Enterprise Simulation within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Enterprise Simulation within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Enterprise Simulation within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- EnterpriseSimulationDefined
- EnterpriseSimulationRegistered
- EnterpriseSimulationStateChanged
- EnterpriseSimulationReviewRequested
- EnterpriseSimulationGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Enterprise Simulation has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.06 - Risk Controls

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-006

Purpose

Define Risk Controls within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Risk Controls within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Risk Controls within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Risk Controls within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.

- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- RiskControlsDefined
- RiskControlsRegistered
- RiskControlsStateChanged
- RiskControlsReviewRequested
- RiskControlsGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.
- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.

- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Risk Controls has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.07 - Executive Control

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-007

Purpose

Define Executive Control within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Executive Control within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Executive Control within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Executive Control within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.

- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- ExecutiveControlDefined
- ExecutiveControlRegistered
- ExecutiveControlStateChanged
- ExecutiveControlReviewRequested
- ExecutiveControlGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.
- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.

- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Executive Control has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.08 - Self Improvement

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-008

Purpose

Define Self Improvement within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Self Improvement within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Self Improvement within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Self Improvement within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.

- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- SelfImprovementDefined
- SelfImprovementRegistered
- SelfImprovementStateChanged
- SelfImprovementReviewRequested
- SelfImprovementGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.
- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.

- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Self Improvement has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.09 - Autonomous Workflows

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-009

Purpose

Define Autonomous Workflows within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomous Workflows within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomous Workflows within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomous Workflows within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomousWorkflowsDefined
- AutonomousWorkflowsRegistered
- AutonomousWorkflowsStateChanged
- AutonomousWorkflowsReviewRequested
- AutonomousWorkflowsGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomous Workflows has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.10 - Autonomous Marketing

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-010

Purpose

Define Autonomous Marketing within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomous Marketing within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomous Marketing within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomous Marketing within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomousMarketingDefined
- AutonomousMarketingRegistered
- AutonomousMarketingStateChanged
- AutonomousMarketingReviewRequested
- AutonomousMarketingGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomous Marketing has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.11 - Autonomous Sales

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-011

Purpose

Define Autonomous Sales within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomous Sales within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomous Sales within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomous Sales within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomousSalesDefined
- AutonomousSalesRegistered
- AutonomousSalesStateChanged
- AutonomousSalesReviewRequested
- AutonomousSalesGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomous Sales has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.12 - Autonomous Delivery

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-012

Purpose

Define Autonomous Delivery within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomous Delivery within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomous Delivery within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomous Delivery within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomousDeliveryDefined
- AutonomousDeliveryRegistered
- AutonomousDeliveryStateChanged
- AutonomousDeliveryReviewRequested
- AutonomousDeliveryGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomous Delivery has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.13 - Client Success Automation

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-013

Purpose

Define Client Success Automation within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Client Success Automation within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation- first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Client Success Automation within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Client Success Automation within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- ClientSuccessAutomationDefined
- ClientSuccessAutomationRegistered
- ClientSuccessAutomationStateChanged
- ClientSuccessAutomationReviewRequested
- ClientSuccessAutomationGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Client Success Automation has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.14 - Growth Automation

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-014

Purpose

Define Growth Automation within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Growth Automation within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Growth Automation within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Growth Automation within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- GrowthAutomationDefined
- GrowthAutomationRegistered
- GrowthAutomationStateChanged
- GrowthAutomationReviewRequested
- GrowthAutomationGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Growth Automation has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.15 - Autonomy Governance

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-015

Purpose

Define Autonomy Governance within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomy Governance within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomy Governance within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomy Governance within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomyGovernanceDefined
- AutonomyGovernanceRegistered
- AutonomyGovernanceStateChanged
- AutonomyGovernanceReviewRequested
- AutonomyGovernanceGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomy Governance has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.16 - Autonomy Events

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-016

Purpose

Define Autonomy Events within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomy Events within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomy Events within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomy Events within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.

- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomyEventsDefined
- AutonomyEventsRegistered
- AutonomyEventsStateChanged
- AutonomyEventsReviewRequested
- AutonomyEventsGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.
- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.

- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomy Events has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.17 - Autonomy Storage

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-017

Purpose

Define Autonomy Storage within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomy Storage within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomy Storage within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomy Storage within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomyStorageDefined
- AutonomyStorageRegistered
- AutonomyStorageStateChanged
- AutonomyStorageReviewRequested
- AutonomyStorageGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomy Storage has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.18 - Autonomy Security

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-018

Purpose

Define Autonomy Security within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomy Security within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomy Security within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomy Security within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomySecurityDefined
- AutonomySecurityRegistered
- AutonomySecurityStateChanged
- AutonomySecurityReviewRequested
- AutonomySecurityGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomy Security has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.19 - Autonomy Testing

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-019

Purpose

Define Autonomy Testing within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomy Testing within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomy Testing within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomy Testing within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomyTestingDefined
- AutonomyTestingRegistered
- AutonomyTestingStateChanged
- AutonomyTestingReviewRequested
- AutonomyTestingGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomy Testing has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.

Book 03.20.20 - Autonomy Roadmap

Version: v22 draft

Status: Draft

Document ID: MAL-V03-B03-20-020

Purpose

Define Autonomy Roadmap within the Autonomous Enterprise blueprint as part of Volume 03 Master Blueprint for Mariam AI Enterprise OS.

Scope

This document covers architecture-level responsibilities, contracts, interfaces, events, storage expectations, security controls, metrics, testing, acceptance criteria, and traceability for Autonomy Roadmap within the Autonomous Enterprise blueprint. It does not implement product code or bypass the documentation-first governance model.

Responsibilities

- Establish a canonical blueprint boundary for Autonomy Roadmap within the Autonomous Enterprise blueprint.
- Convert strategic roadmap intent into implementation-ready subsystem contracts.
- Coordinate with Enterprise Core identity, runtime, event, storage, security, and observability services.
- Preserve governance, traceability, reviewability, and rollback expectations for high-impact behavior.
- Provide enough detail for future specifications, testing guides, operations guides, and implementation issues.

Inputs

- Volume 00 Enterprise Constitution principles for governance, security, quality, engineering, AI, DNA, swarm, and evolution.
- Volume 01 Master Vision strategy for enterprise outcomes, knowledge leverage, marketplace expansion, and autonomous work.
- Volume 02 Master Roadmap phase guidance corresponding to Autonomous Enterprise.
- Earlier Volume 03 blueprint books for Enterprise Core, Enterprise Organization, AI Society, Knowledge System, and Capability System where dependencies apply.
- Domain requirements, user intent, policy context, runtime constraints, and release acceptance evidence.

Outputs

- Blueprint contract and subsystem boundaries for Autonomy Roadmap within the Autonomous Enterprise blueprint.
- Canonical records, events, metrics, storage expectations, and security controls.
- Traceable inputs for Volume 06 specifications and later specialist volumes.
- Acceptance evidence that downstream implementation can use before writing system code.

Interfaces

- Enterprise Core kernel, runtime manager, event bus, object manager, identity access, storage, security, and observability interfaces.
- Enterprise Organization decision, approval, escalation, KPI, and operating rhythm interfaces.
- AI Society agent, planning, execution, review, validation, and governance interfaces.
- Knowledge System graph, store, search, indexing, provenance, and quality interfaces.
- Capability System registry, matching, lifecycle, benchmark, governance, and execution interfaces.

Events

- AutonomyRoadmapDefined
- AutonomyRoadmapRegistered
- AutonomyRoadmapStateChanged
- AutonomyRoadmapReviewRequested
- AutonomyRoadmapGovernanceBlocked

Storage

- Canonical records for definitions, versions, owners, state, policy classification, and dependency metadata.
- Operational journals for lifecycle changes, execution evidence, review outcomes, exceptions, and acceptance decisions.
- Audit metadata linking user intent, agent or service identity, source context, runtime, model or provider choice, and policy status.
- Retention metadata that separates draft, active, archived, deprecated, and superseded records.

Security

- Apply least privilege to users, agents, services, providers, plugins, connectors, runtimes, and storage records.
- Deny protected actions when required permissions, policy checks, evidence, or human review gates are missing.
- Classify inputs, outputs, events, logs, metrics, memory, and exported artifacts before sharing or persistence.
- Audit privileged actions, denied actions, governance exceptions, configuration changes, and lifecycle transitions.
- Prevent generated or inferred outputs from becoming authoritative without explicit acceptance rules.

Metrics

- Adoption rate, successful completion rate, rejection rate, rework count, and user trust indicators.
- Latency, cost, throughput, availability, error rate, fallback rate, and recovery time.
- Quality score, benchmark pass rate, validation pass rate, review backlog, and stale-record count.
- Security denial count, governance exception count, policy violation count, and audit completeness.

Testing

- Contract tests for required inputs, outputs, interfaces, event envelopes, storage records, and lifecycle states.
- Security tests for tenant isolation, least privilege, protected action denial, redaction, and audit integrity.
- Quality tests for correctness, repeatability, ranking or routing behavior, recovery paths, and acceptance evidence.

- Integration tests across Enterprise Core, Organization, AI Society, Knowledge System, and Capability System dependencies.
- Release tests verifying Markdown inclusion, PDF generation, ZIP packaging, manifest checksums, and release notes.

Acceptance Criteria

- Autonomy Roadmap has explicit ownership, lifecycle, interfaces, events, storage, security, metrics, and testing expectations.
- The blueprint is detailed enough for downstream specifications without requiring implementation code in this repository.
- Governance, security, and review controls are visible and traceable.
- Release artifacts include this Markdown file, the book PDF, MANIFEST.json, release notes, and the versioned ZIP.

Traceability

- MAL-V00-B009
- MAL-V00-B010
- MAL-V00-B011
- MAL-V00-B012
- MAL-V02-B020
- MAL-V03-B03-01-004
- MAL-V03-B03-05-003

Version History

Version	Date	Status	Notes
v22 draft	2026-06-29	Draft	Initial Autonomous Enterprise blueprint content for Mariam Architecture Library v22.